

Mining Frequent Subgraphs

Prof. O-Joun Lee

Dept. of Artificial Intelligence,
The Catholic University of Korea
ojlee@catholic.ac.kr

Contents



- Frequent subgraph mining (FSM)
- FSM algorithms
 - Apriori-Based Approach: FSG
 - DFS Approach
 - Subdue Approach
- Sample code: Mining frequent subgraphs in a graph using the gSpan algorithm in NetworkX


- Frequent pattern: a structure (a set of items, subsequences, substructures, etc.) that occurs frequently in a data set
- Motivation: Finding inherent regularities in data
 - What products were often purchased together?
 - What are the subsequent purchases after buying a PC?
 - What sequences of DNA are sensitive to this new drug?
 - Which topics are in a collection of documents?


- Also called Downward closure Property
- All subsets of a frequent pattern must also be frequent
 - Because any item that contains X must also contains subset of X .

If we have already verified that X is infrequent, there is no need to count X 's supersets because they **MUST** be infrequent too.

- Frequent subgraphs
 - A (sub)graph is frequent if its support (occurrence frequency) in a given dataset is no less than a minimum support threshold
- Applications of graph pattern mining:
 - Mining biochemical structures
 - Program control flow analysis
 - Mining XML structures or Web communities
 - Building blocks for graph classification, clustering, compression, comparison, and correlation analysis

- Problem: Find all subgraphs of G that appear at least t times
- Suppose $t = 2$, the frequent subgraphs are (only edge labels)
 - a, b, c
 - $a-a, a-c, b-c, c-c$
 - $a-c-a \dots$
- Exponential number of patterns!






Support	1	3	3
Subgraph			


- Apriori-based approaches
 - Start with small-size subgraphs and proceeds in a bottom-up manner
 - Join two patterns to create bigger size patterns (through Apriori principle)
 - Several approaches
 - FSG
 - PATH#
- Pattern-growth approaches
 - Extends existing frequent graphs by adding one edge
 - Several approaches:
 - gSpan, MoFa
 - Gaston FFSM, SPIN
- Greedy approaches
 - Subdue

- Start with small-size subgraphs and proceeds in a bottom-up manner
- Join two patterns to create bigger size patterns (through Apriori principle)



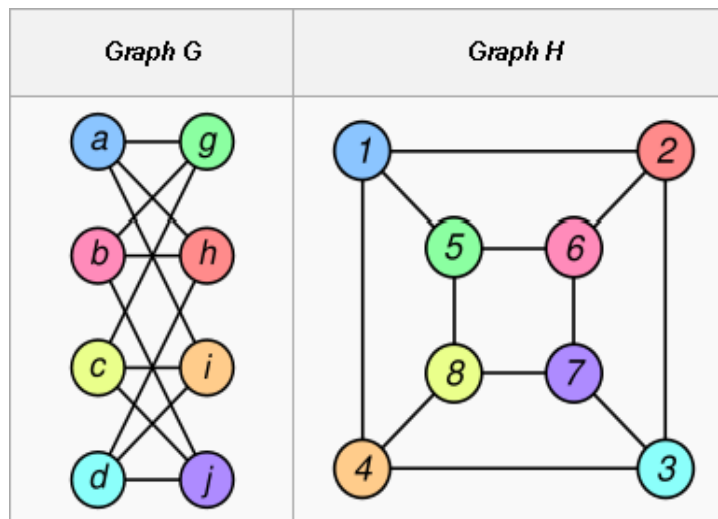
- Problem:
 - Join operation among graphs is extremely expensive

- Generate patterns expanding existing ones
- Extends existing frequent graphs by adding one edge



- Problems
 - Duplicate graphs

- Graph isomorphism
 - To detect if two graphs are identical in structure
- Graph representation (Canonical Labeling)
 - A canonical label is a **unique code** of a given graph.
 - Canonical label should be **the same no matter how graphs are represented**, as long as graphs have the same topological structure and the same labeling of edges and vertices.
- Subgraph candidate generation
 - Generate candidate frequent subgraphs from datasets



- Returns True if the graphs G1 and G2 are isomorphic and False otherwise.
(The two graphs G1 and G2 must be the same type.)

```
#Import networkx, isomorphism
import networkx as nx
import networkx.algorithms.isomorphism as iso

#create graph 1
G1 = nx.Graph()

G1.add_nodes_from(['A','B','C','D','E','F'])

G1.add_edges_from([('A','B'),('A','C'),('A','D'),('A','E'),('A','F')])

#create graph 2
G2 = nx.star_graph(5)
```

Testing if two graphs are isomorphic

```
nx.is_isomorphic(G1, G2)
```

True

➤ How to find edge mapping

```
import networkx as nx
G1 = nx.Graph()
G1.add_weighted_edges_from([(0,1,0), (0,2,1), (0,3,2)], weight = 'aardvark')
G2 = nx.Graph()
G2.add_weighted_edges_from([(0,1,0), (0,2,2), (0,3,1)], weight = 'baboon')
G3 = nx.Graph()
G3.add_weighted_edges_from([(0,1,0), (0,2,2), (0,3,2)], weight = 'baboon')

def comparison(D1, D2):
    #for an edge u,v in first graph and x,y in second graph
    #this tests if the attribute 'aardvark' of edge u,v is the
    #same as the attribute 'baboon' of edge x,y.

    return D1['aardvark'] == D2['baboon']


nx.is_isomorphic(G1, G2, edge_match = comparison)
```

True


```
nx.is_isomorphic(G1, G3, edge_match = comparison)
```

False

- Given a set of 4 graphs:



- Support: frequency of a subgraph appearing in a set of graphs
- Frequent subgraph Min support = $3/4$



Apriori principle (for graphs): If a graph is frequent, all of its subgraphs are frequent

- We define a labeled graph G as a five element tuple $G = \{V, E, \Sigma_V, \Sigma_E, \delta\}$




Where:

V is the set of vertices of G ,

$E \subseteq V \times V$ is a set of undirected edges of G ,

Σ_V (Σ_E) are set of vertex (edge) labels,




δ is the labeling function: $V \rightarrow \Sigma_V$ and $E \rightarrow \Sigma_E$ that maps vertices and edges to their labels.



- Support: Given a set of labeled graphs:
 - $D = \{G_1, G_2, \dots, G_n\}, G_i = \langle V_i, E_i, \ell_i \rangle$
 - A subgraph G .
- The supporting set of G is: $D_G = \{G_i | G \sqsubseteq G_i, G_i \in D\}$
Where $G \sqsubseteq G_i$ indicates that G is subgraph isomorphic to G_i
- The support is defined as:

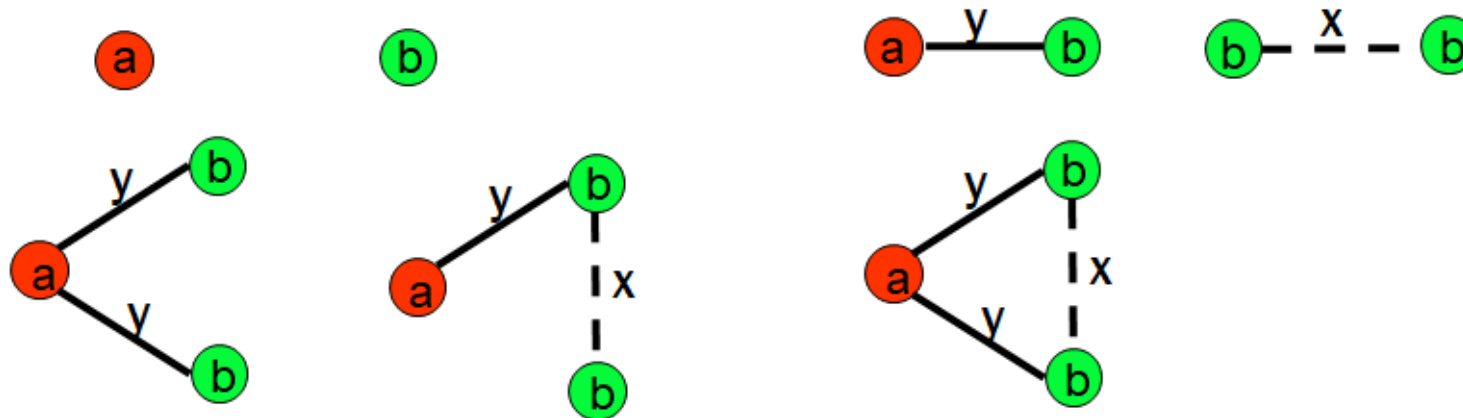
$$\sigma(G) = \frac{|D_G|}{|D|}$$

- *Input: A set Graph data of labeled undirected graphs*



$$\sigma = 2/3$$

- All frequent subgraphs (w. r. t. σ) from *Graph data*.




- Input
 - Set of labeled-graphs $D = \{G_1, G_2, \dots, G_n\}$, $G_i = \langle V_i, E_i, \ell_i \rangle$
 - Minimum support min_sup
- Output:
 - A subgraph G is frequent if $\sigma(G) \geq \text{min_sup}$
 - Each subgraph is connected.


- Apriori-based approaches:
 - FSG
- Pattern-growth approaches:
 - gSpan
- Greedy approach:
 - Subdue

Methodology: breadth-search, joining two graphs

- 1. Generates new graphs with **one more node**



- 2. Generates new graphs with **one more edge**




- $K = 1$
- $F_1 =$ all frequent edges
- Repeat
 - $K = K + 1$;
 - $C_K = \text{join}(F_{K-1})$
 - $F_K =$ frequent patterns in C_K
 - Until F_K is empty


- $K = 1$
- $F_1 =$ all frequent edges
- Repeat
 - $K = K + 1$;
 - $C_K = \text{join}(F_{K-1})$
 - $F_K =$ frequent patterns in C_K
 - Until F_K is empty

- $\text{Join}(L) = \cup \text{join}(P, Q)$ for all $P, Q \in L$
- $\text{Join}(P, Q) = \{G \mid P, Q, \subset G, |G| = |P| + 1, |P| = |Q|\}$
- Two graphs P and Q are **joinable** if the join of the two graphs produces a non-empty set
- Theorem: two graphs P and Q are joinable if $P \cap Q$ is a graph with size $|P| - 1$ or share a common “core” with size $|P| - 1$




Case 1: identical node labels




Case 2: Core contains identical labels








Core: The $(k-1)$ subgraph that is common between the joint graphs



Case 3: Core multiplicity





+

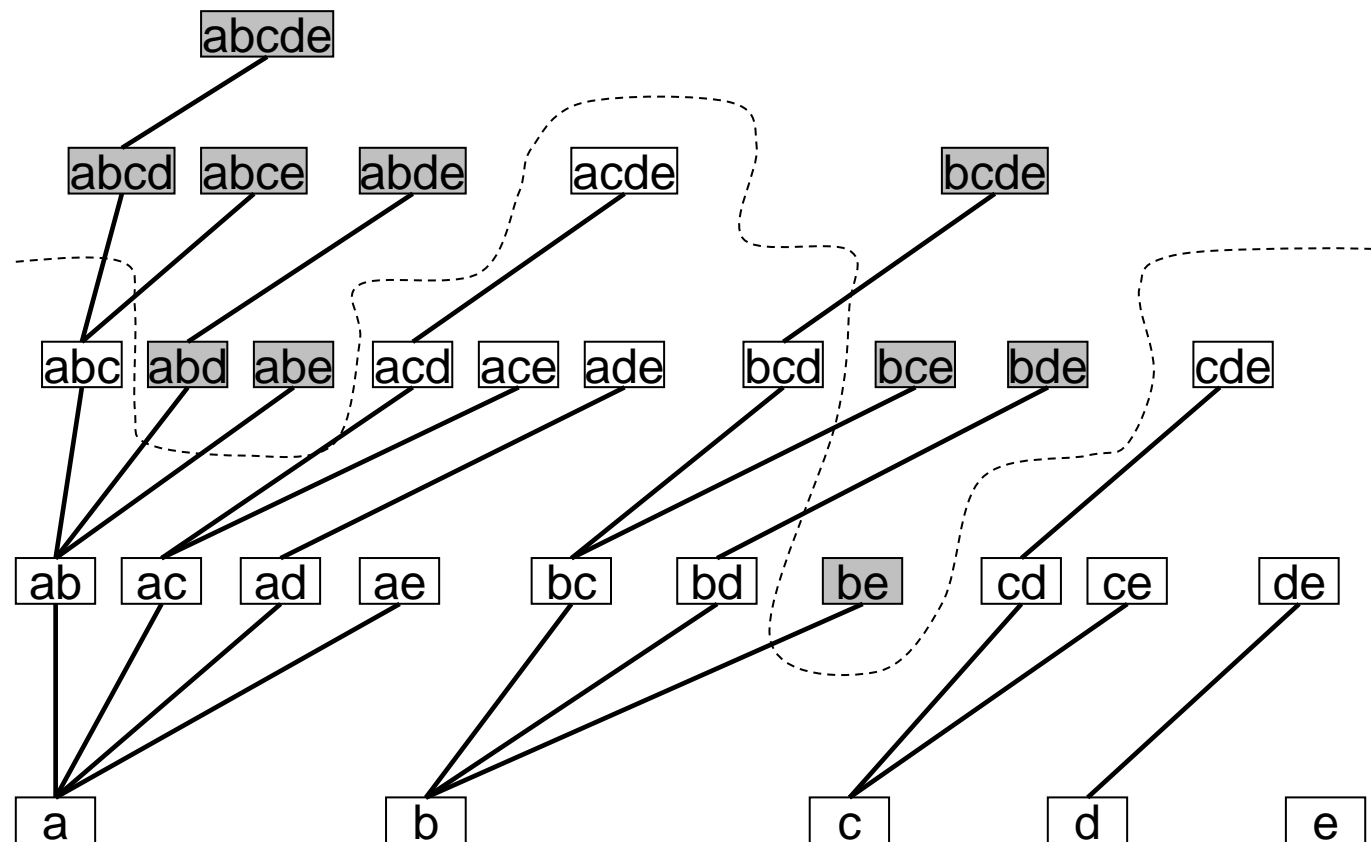


- Graph isomorphism
 - Two graphs may have the same topology though their layouts are different
- Subgraph isomorphism
 - How to compute the support value of a pattern

- A graph is isomorphic if it is topologically equivalent to another graph




➤ Itemset search space – prefix based




- Canonical representation of itemset is obtained by a complete order over the items.
- Each possible itemset appear in TSS exactly once - no duplications or omissions.
- Properties of Tree search space
 - for each k-label, its parent is the k-1 prefix of the given k-label
 - The relation among siblings is in ascending lexicographic order.



- Organize DFS code nodes as parent-child.
- Pre-order traversal follows DFS lexicographic order.
- If s and s' are the same graph with different DFS codes, s' is not the minimum and can be pruned.




- Given a graph G . for each Depth First Search over graph G , construct the corresponding DFS-Code.




	(a)	(b)	(c)
1	(0, 1, X, a, Y)	(0, 1, Y, a, X)	(0, 1, X, a, X)
2	(1, 2, Y, b, X)	(1, 2, X, a, X)	(1, 2, X, a, Y)
3	(2, 0, X, a, X)	(2, 0, X, b, Y)	(2, 0, Y, b, X)
4	(2, 3, X, c, Z)	(2, 3, X, c, Z)	(2, 3, Y, b, Z)
5	(3, 1, Z, b, Y)	(3, 0, Z, b, Y)	(3, 0, Z, c, X)
6	(1, 4, Y, d, Z)	(0, 4, Y, d, Z)	(2, 4, Y, d, Z)



(a)





(b)




(c)

Min
DFS-Code


	(a)	(b)	(c)
1	(0, 1, X, a, Y)	(0, 1, Y, a, X)	(0, 1, X, a, X)
2	(1, 2, Y, b, X)	(1, 2, X, a, X)	(1, 2, X, a, Y)
3	(2, 0, X, a, X)	(2, 0, X, b, Y)	(2, 0, Y, b, X)
4	(2, 3, X, c, Z)	(2, 3, X, c, Z)	(2, 3, Y, b, Z)
5	(3, 1, Z, b, Y)	(3, 0, Z, b, Y)	(3, 0, Z, c, X)
6	(1, 4, Y, d, Z)	(0, 4, Y, d, Z)	(2, 4, Y, d, Z)



(a)



(b)




(c)

- The minimum DFS code $\min(G)$, in DFS lexicographic order, is a canonical representation of graph G .
- Graphs A and B are isomorphic if and only if:

$$\min(A) = \min(B)$$

- If $\min(G_1) = \{a_0, a_1, \dots, a_n\}$
and $\min(G_2) = \{a_0, a_1, \dots, a_n, b\}$
 G_1 is parent of G_2
 G_2 is child of G_1
- A valid DFS code requires that **b** grows from a vertex on the rightmost path (inherited property from the DFS search).




$\text{Min}(g) = (0,1,X,a,Y) (1,2,Y,b,X) (2,0,X,a,X) (2,3,X,c,Z) (3,1,Z,b,Y) (1,4,Y,d,Z)$


A child of graph G_1 *must grow edge from rightmost path of G_1* (necessary condition)

Graph G_2 :

wrong

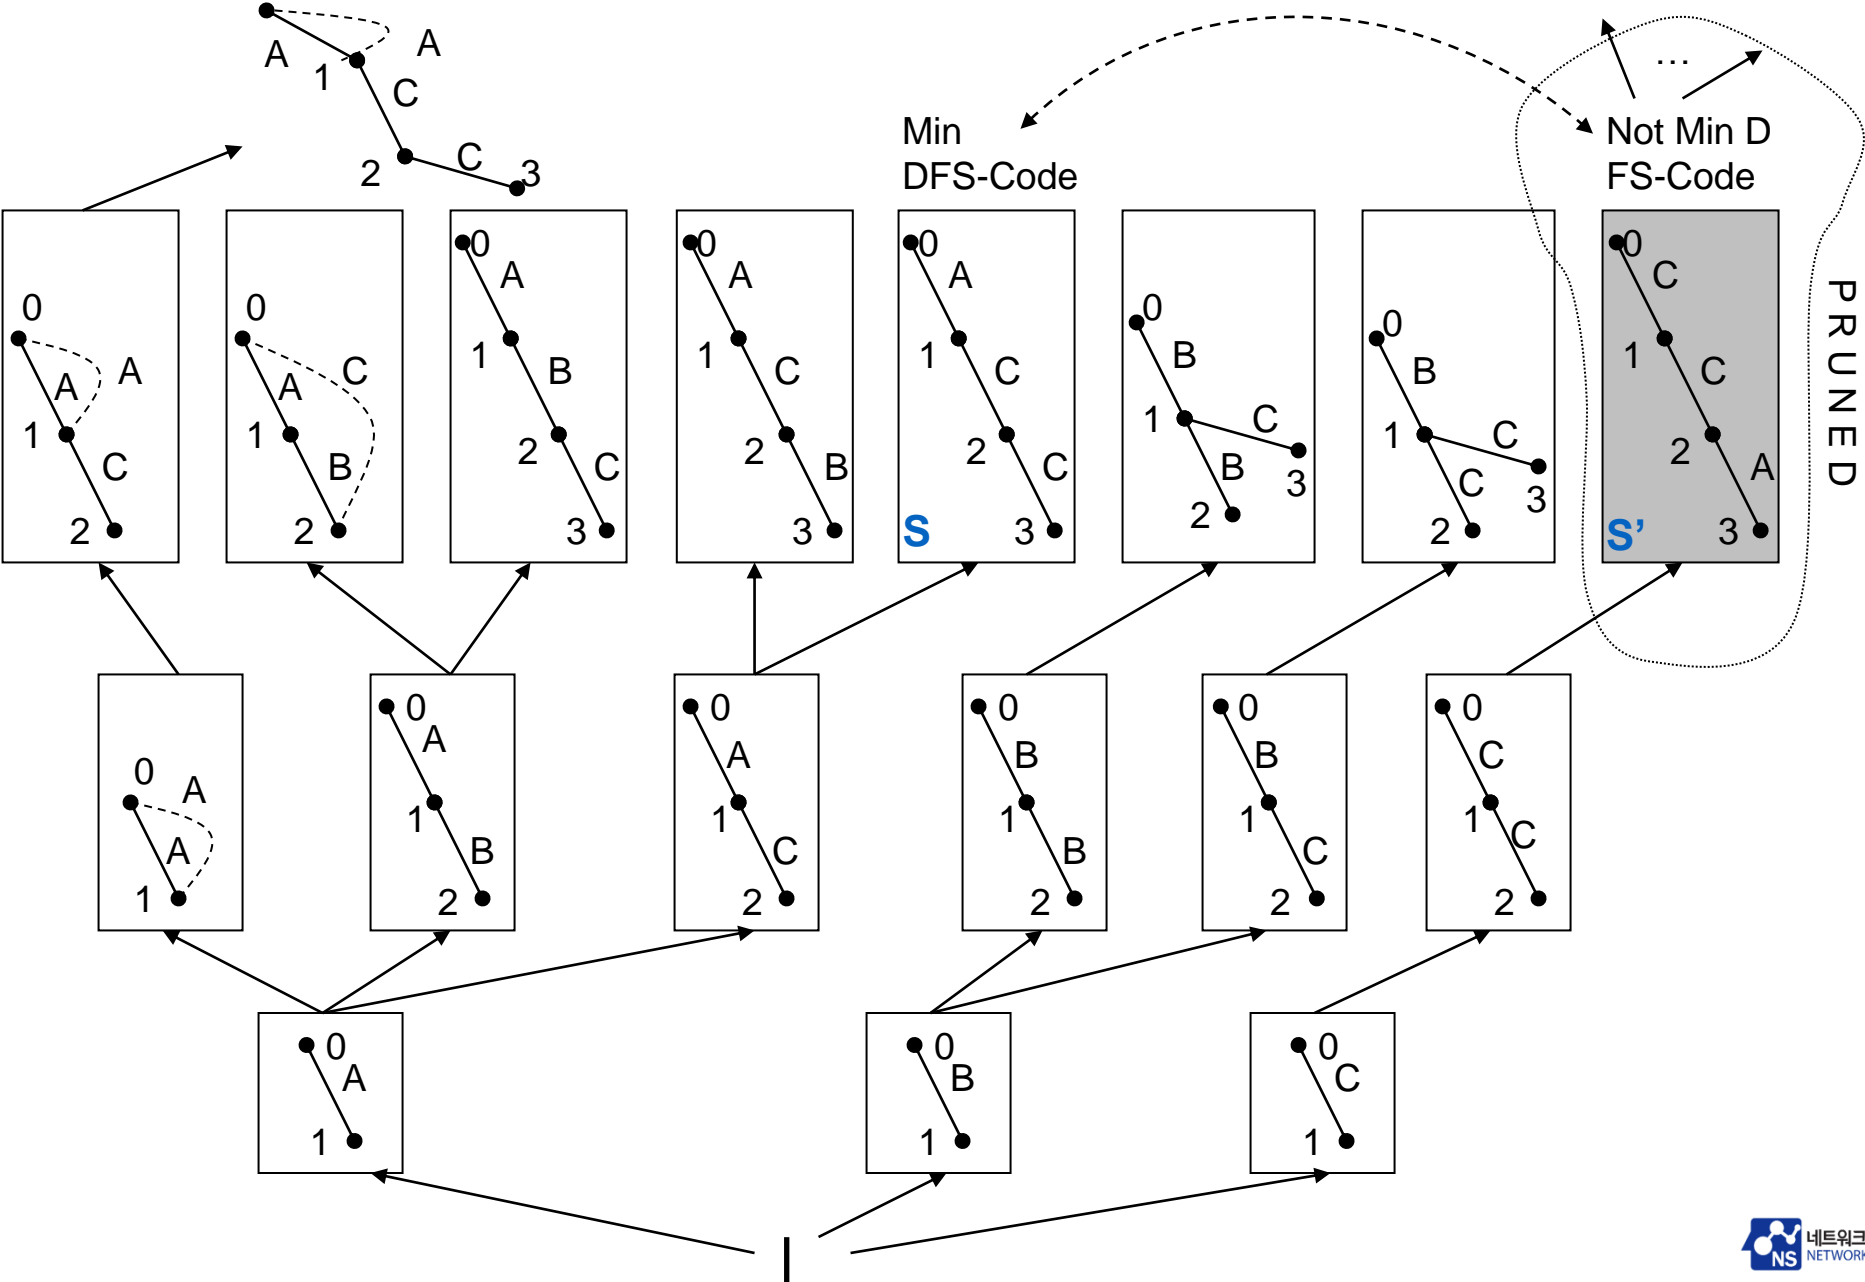


Forward edge



Backward edge

- Organize DFS Code nodes as parent-child.
- Sibling nodes organized in ascending DFS lexicographic order.
- *InOrder* traversal follows DFS lexicographic order



- All of the descendants of infrequent node are infrequent also.
- All of the descendants of a not minimal DFS code are also not minimal DFS codes.

gSpan(D, F, g)

1: if $g \neq \min(g)$

 return;

2: $F \leftarrow F \cup \{ g \}$

3: $\text{children}(g) \leftarrow$ [generate all g' potential children with one edge growth]




4: Enumerate($D, g, \text{children}(g)$)

5: for each $c \in \text{children}(g)$


 if $\text{support}(c) \geq \# \text{minSup}$

 SubgraphMining (D, F, c)

Given: database D



Task: Mine all frequent subgraphs with support ≥ 2 (#minSup)



➤ Support: 2








```
from gspan_mining.config import parser
from gspan_mining.main import main
```

```
%pylab inline
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.
Populating the interactive namespace from numpy and matplotlib

```
args_str = '-s 2 -l 3 -p True ./graphdata/sample_data3'
FLAGS, _ = parser.parse_known_args(args=args_str.split())
```

```
gs = main(FLAGS)
print(gs)
```

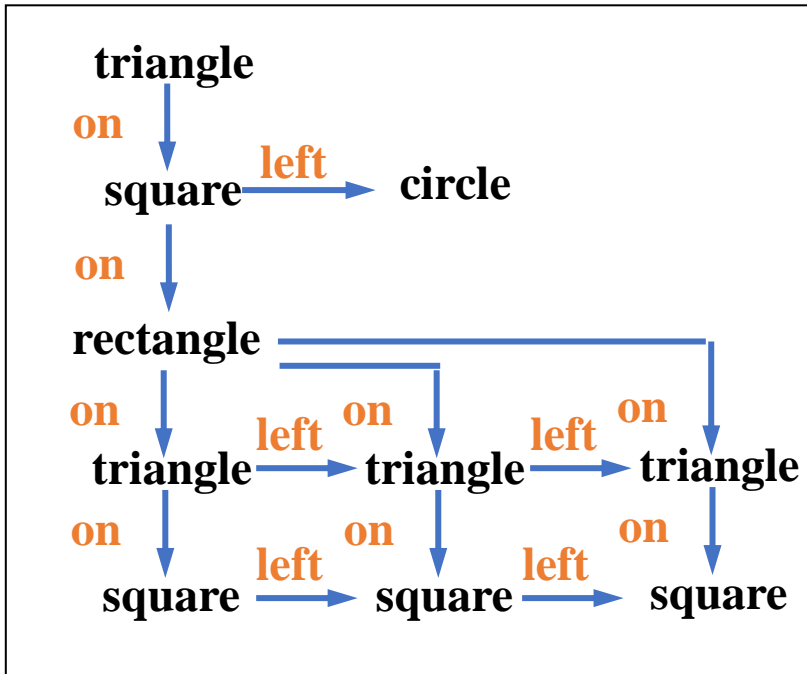


- A greedy algorithm for finding some of the most prevalent subgraphs.
- This method is not complete, i.e. it may not obtain all frequent subgraphs, although it pays in fast execution.

- It discovers substructures that compress the original data and represent structural concepts in the data.
- Based on *Beam Search* - like BFS it progresses level by level. Unlike BFS, however, beam search moves downward only through the best W nodes at each level. The other nodes are ignored.

- Step 1: Create substructure for each unique vertex label

DB:

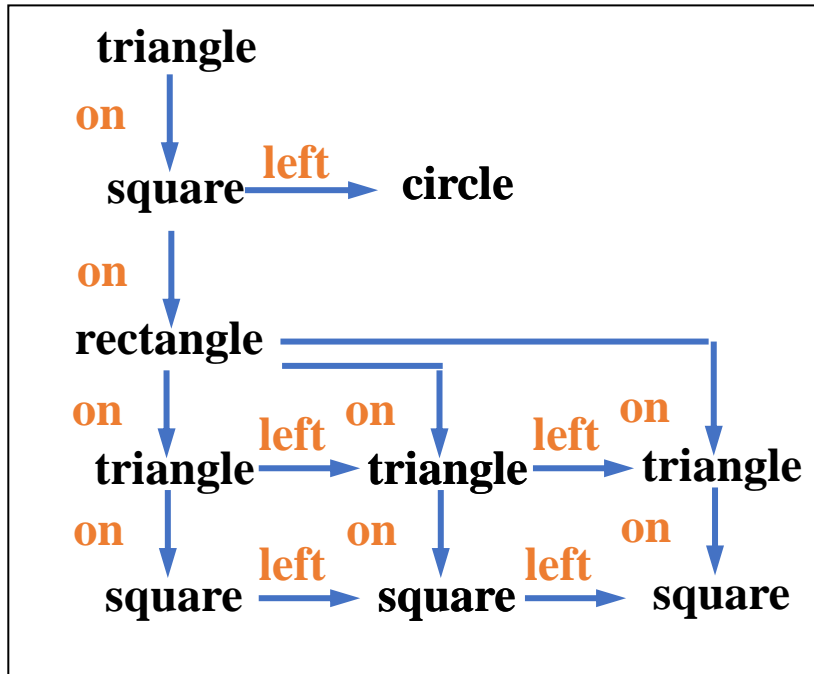


Substructures:

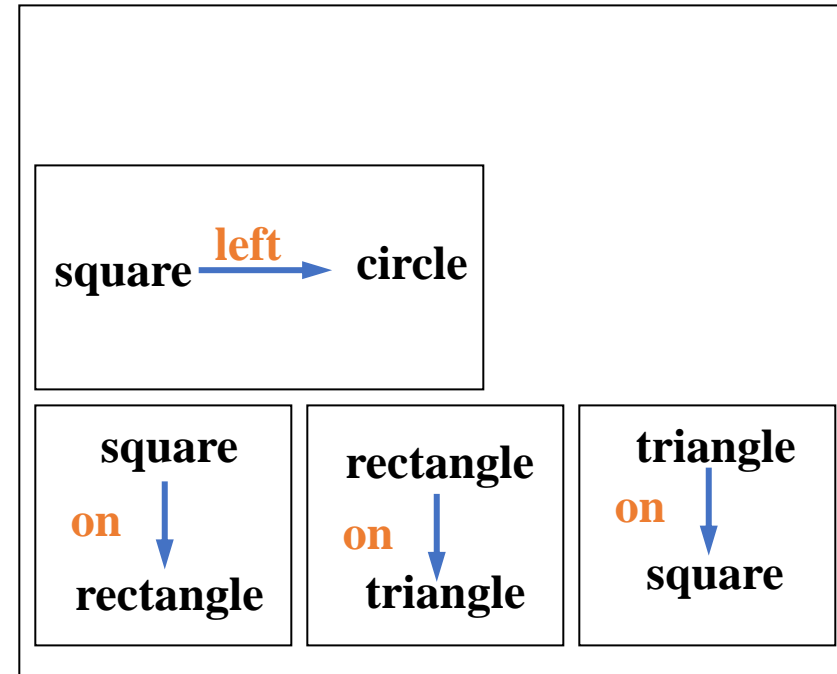
triangle (4)
square (4)
circle (1)
rectangle (1)

- Step 2: Expand best substructure by an edge or edge and neighboring vertex

DB:



Substructures:



- Step 3: Keep only best substructures on queue (specified by beam width).
- Step 4: Terminate when queue is empty or when the number of discovered substructures is greater than or equal to the limit specified.
- Step 5: Compress graph and repeat to generate hierarchical description.



네트워크 과학연구실
NETWORK SCIENCE LAB



가톨릭대학교
THE CATHOLIC UNIVERSITY OF KOREA

