

Node Classification

Prof. O-Joun Lee

Dept. of Artificial Intelligence,
The Catholic University of Korea
ojlee@catholic.ac.kr



네트워크 과학연구실
NETWORK SCIENCE LAB




가톨릭대학교
THE CATHOLIC UNIVERSITY OF KOREA

Contents



- Overview of node classification
- Feature extraction methods
- Node classification approaches
- Evaluation metrics

- Given a graph and few nodes for which we know the “**label**” or a “**class**,” how can we predict user attributes or interests?




Predict the labels for non marked nodes?

- Is this a friend or an acquaintance?
- **Recommendation** systems to suggest objects (music, movies, activities)
- Automatically **understand roles in a network** (hubs, activators, influencing nodes, etc.)
- Identify experts for question answering systems
- Targeted advertising
- Study of communities (key individuals, group starters ...)
- Study of diseases and cures
- Identify **unusual behaviours** or behavioural changes
- Finding similar nodes and outliers

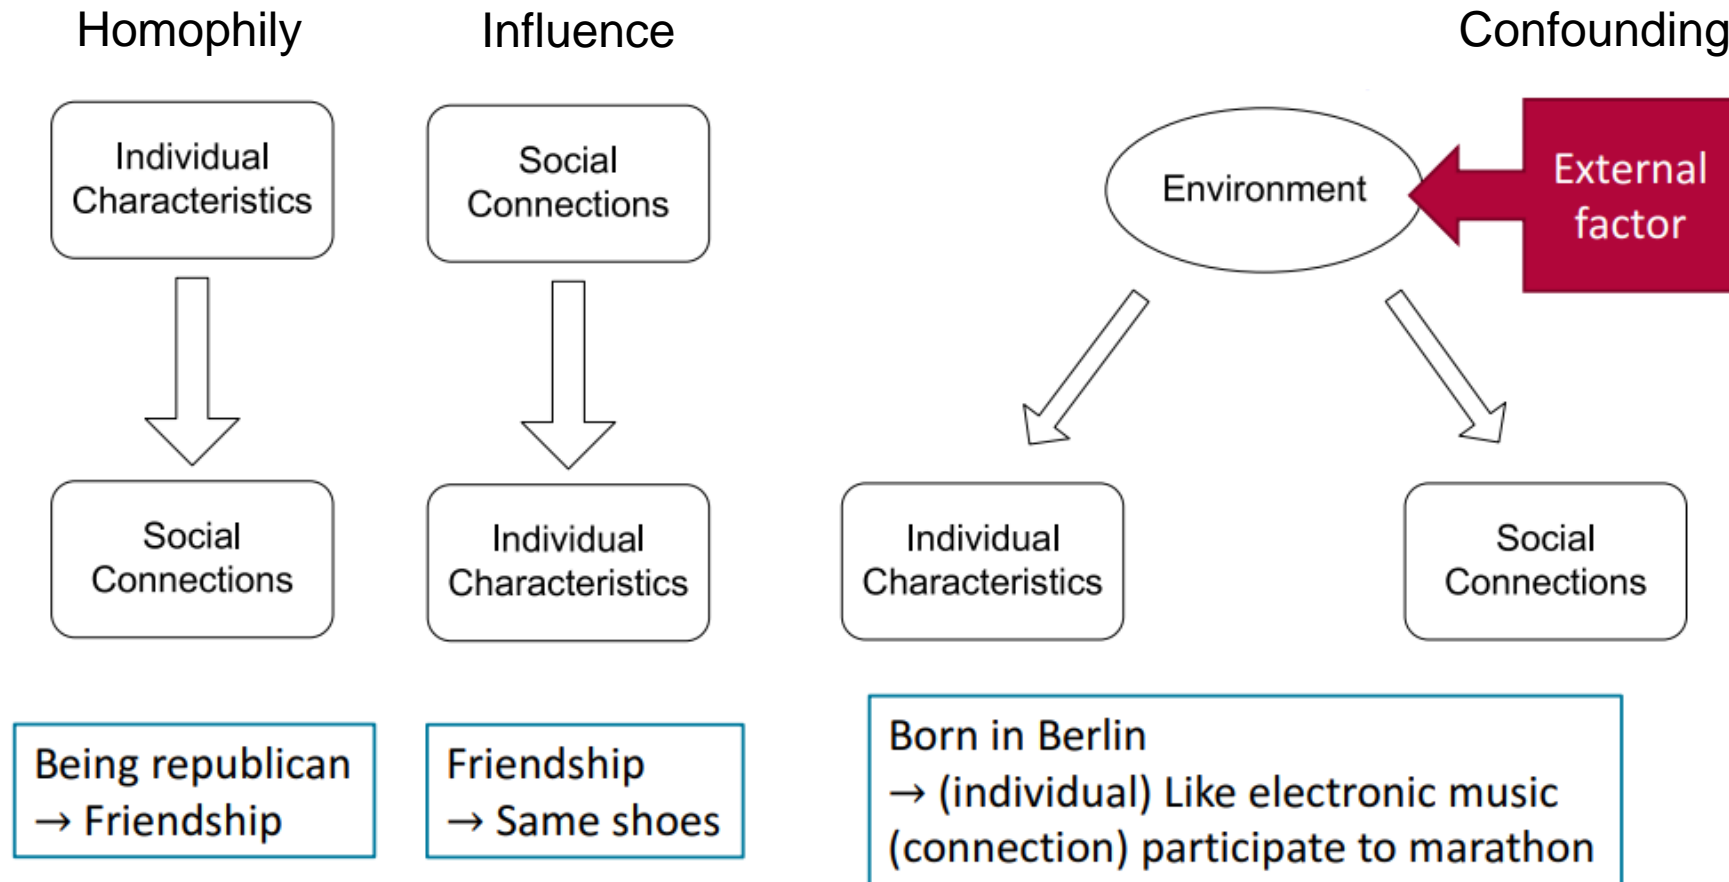
- **Not all the nodes have labels**
 - users are not willing to provide explanations
- Roles are **not explicitly** declared
 - who is more important in a company? (Think about the exchanged emails)
- Labels provided by the users can be misleading
- Labels are sparse
 - some categories might be missing or incomplete

- Given:
 - Graph $G = (V, E, W)$ with vertices V , edges E and weight matrix W
 - Labeled nodes $V_{lab} \subset V$, unlabeled nodes $V_{un} = V \setminus V_{lab}$
 - Y the set of m possible labels (e.g., $Y = \{\text{republican}, \text{democrat}\}$)
 - $Y_{lab} = \{y_1, y_2, \dots, y_l\}$ the labels on labeled nodes in V_{lab}
- Problem:
 - Infer labels Y_{un} for all nodes in V_{un}




- Can be generalized to multilabel and multiclass classification:
 - With multiclass classification assume that each labelled node has a **probability distribution on the labels**.
- Can work on generalized graph structures
 - hypergraphs, graphs with weighted, labelled, timestamped edges, multigraphs, probabilistic graphs and so on.

- Individual **behaviours** are correlated in a network **environment**



- The **graph structure** encodes important **information for node classification**
- So, it is reasonable to think that:
 - **labels propagate** in the network **following the links**
- Methods that work with points in the space perform poorly in a graph




Assumption: The label propagates on the network

- Node features:
 - Measurable characteristics of the nodes that help
 - discriminating a node from another
 - or stating the similarity with other nodes.
- Examples of features:
 - In/out degree of the node
 - Number of L-labelled edges from that node
 - Number of paths in that goes through the node
 - Number of triangles
 - Degree and number within ego-net edges
 - etc.

- **Similarity based**
 - Find nodes that **share the same characteristics** with other nodes
- **Iterative learning**
 - Learn a set of labels and **propagate the information to similar nodes**
- **Label propagation**
 - Labelled nodes **propagate the information to the neighbours** with some probability

➤ Real-world Applications:

Customers Who Bought This Item Also Bought




An Introduction to Statistical Learning:...
› Gareth James
★★★★★ 41
#1 Best Seller in Mathematical & Statistical...
Hardcover
\$56.75 Prime




Advanced R (Chapman & Hall/CRC The R...
› Hadley Wickham
★★★★★ 13
Paperback
\$54.05 Prime



Statistical Analysis of Network Data with R...
› Eric D. Kolaczyk
★★★★☆ 1
Paperback
\$53.93 Prime




Machine Learning with R
Brett Lantz
★★★★☆ 28
Paperback
\$49.49 Prime



Scholar articles

On power-law relationships of the internet topology
M Faloutsos, P Faloutsos, C Faloutsos - ACM SIGCOMM Computer Communication Review, 1999
Cited by 5151 - **Related articles** - All 88 versions



➤ Movies recommendations



The Social Network
has been added to the
Saved section of your DVD
Queue

Availability:
DVD: Unknown

More like The Social Network

Zodiac



Add

☆☆☆☆☆

Not Interested

'N Sync: The Reel 'N Sync



Add

☆☆☆☆☆

Not Interested

Edison Force




Add

☆☆☆☆☆

Not Interested

네트워크 과학 연구실 NETWORK SCIENCE LAB 가톨릭대학교 THE CATHOLIC UNIVERSITY OF KOREA

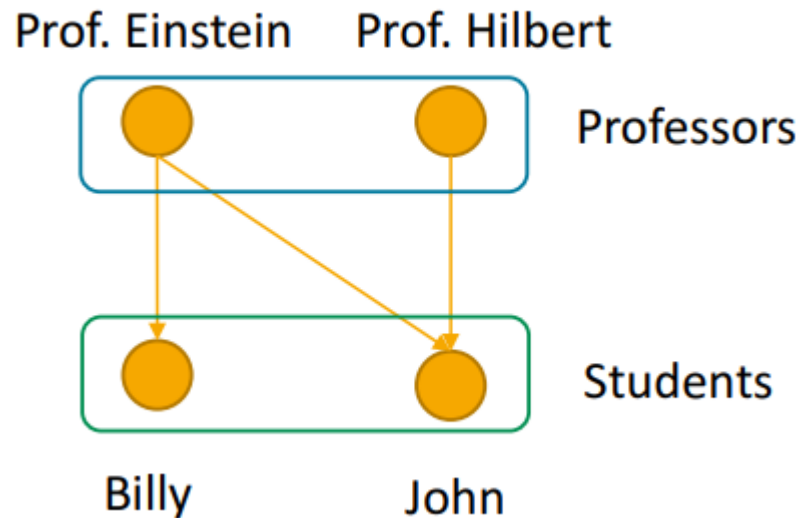
- Topical search engine is an engine that focuses on a particular topic.
- It covers a part of the whole Web rather than a particular website - this is possible because Programmable Search Engine allows you to include multiple websites in the same engine.



- **Equivalences in terms of structure**
 - Structural, Automorphic, and Regular
- **Role** extraction methods:
 - RoIX
- **Recursive** similarities
 - Paths, Max-flow, **SimRank**

- Two nodes u and v are **structurally equivalent** if they **have the same relationships to all other nodes**.
- Two nodes u and v are **automorphically equivalent** if all the nodes can be **relabelled to form an isomorphic graph** with the labels of u and v interchanged (just change the node id).
- Two nodes u and v are **regularly equivalent** if they are **equally related to equivalent others**


- Two nodes u and v are regularly equivalent if they are **equally related to equivalent others**
- Assumes a similarity between sets of nodes



Billy and John are similar because they are **both connected to a professor**.
Same for prof. Einstein and Hilbert.

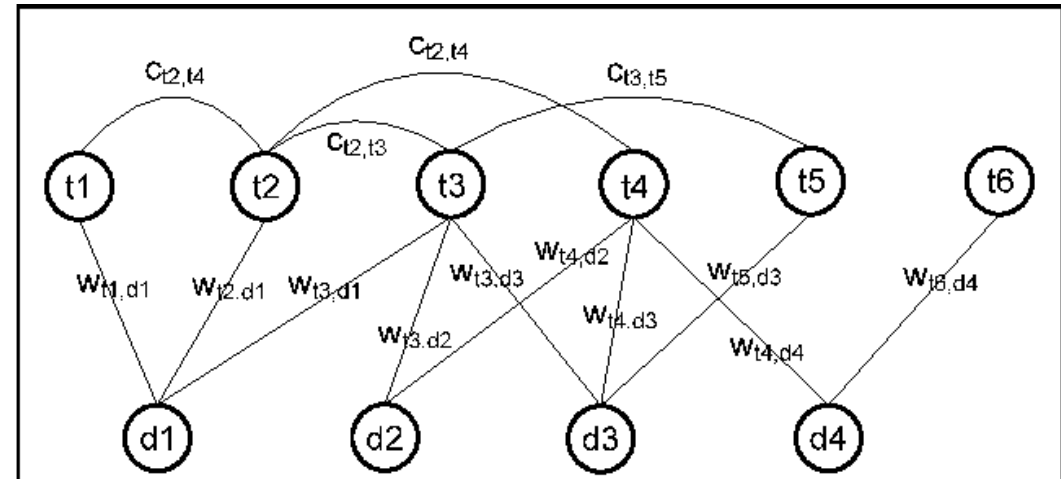
Regular equivalence doesn't care about which connections but to **which set/group a node is connected**

- Two nodes u and v are regularly equivalent if they are **equally related to equivalent others**
- Assumes a similarity between sets of nodes

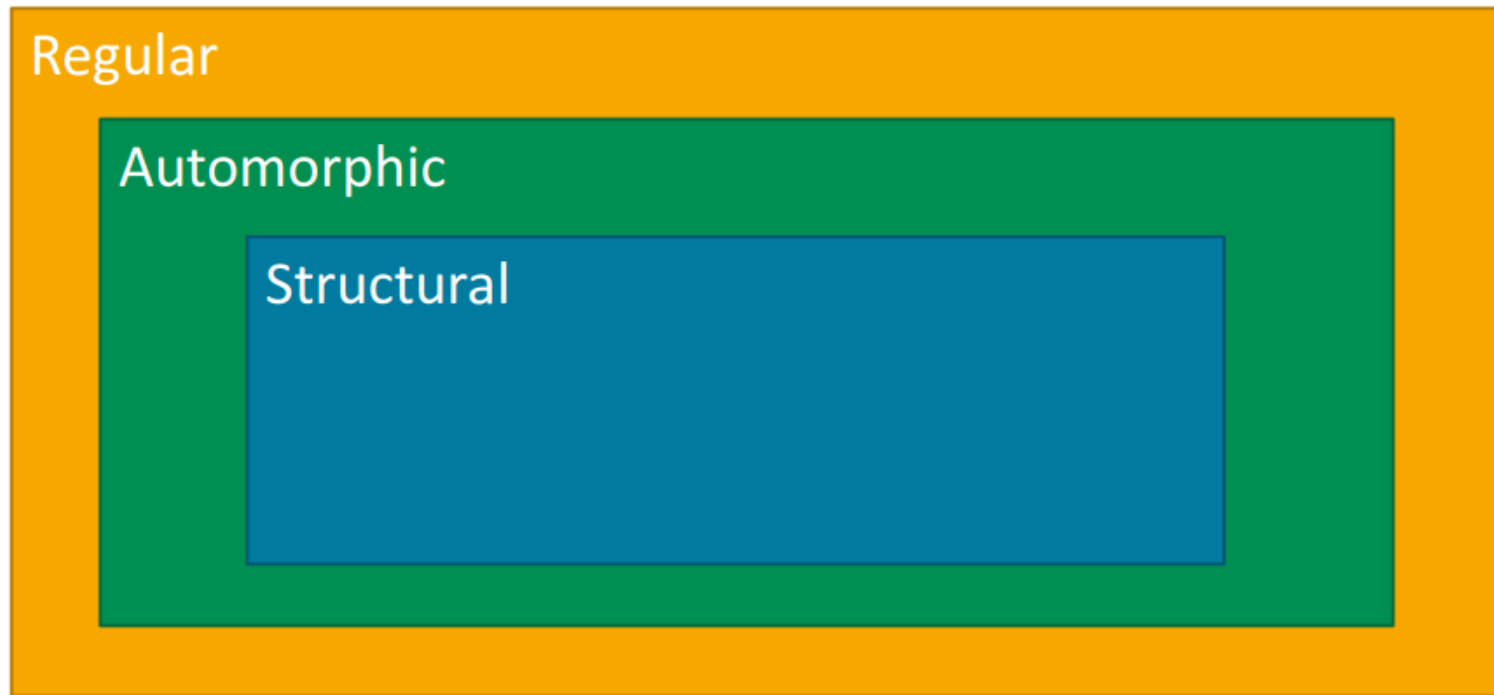


—————→ Interact

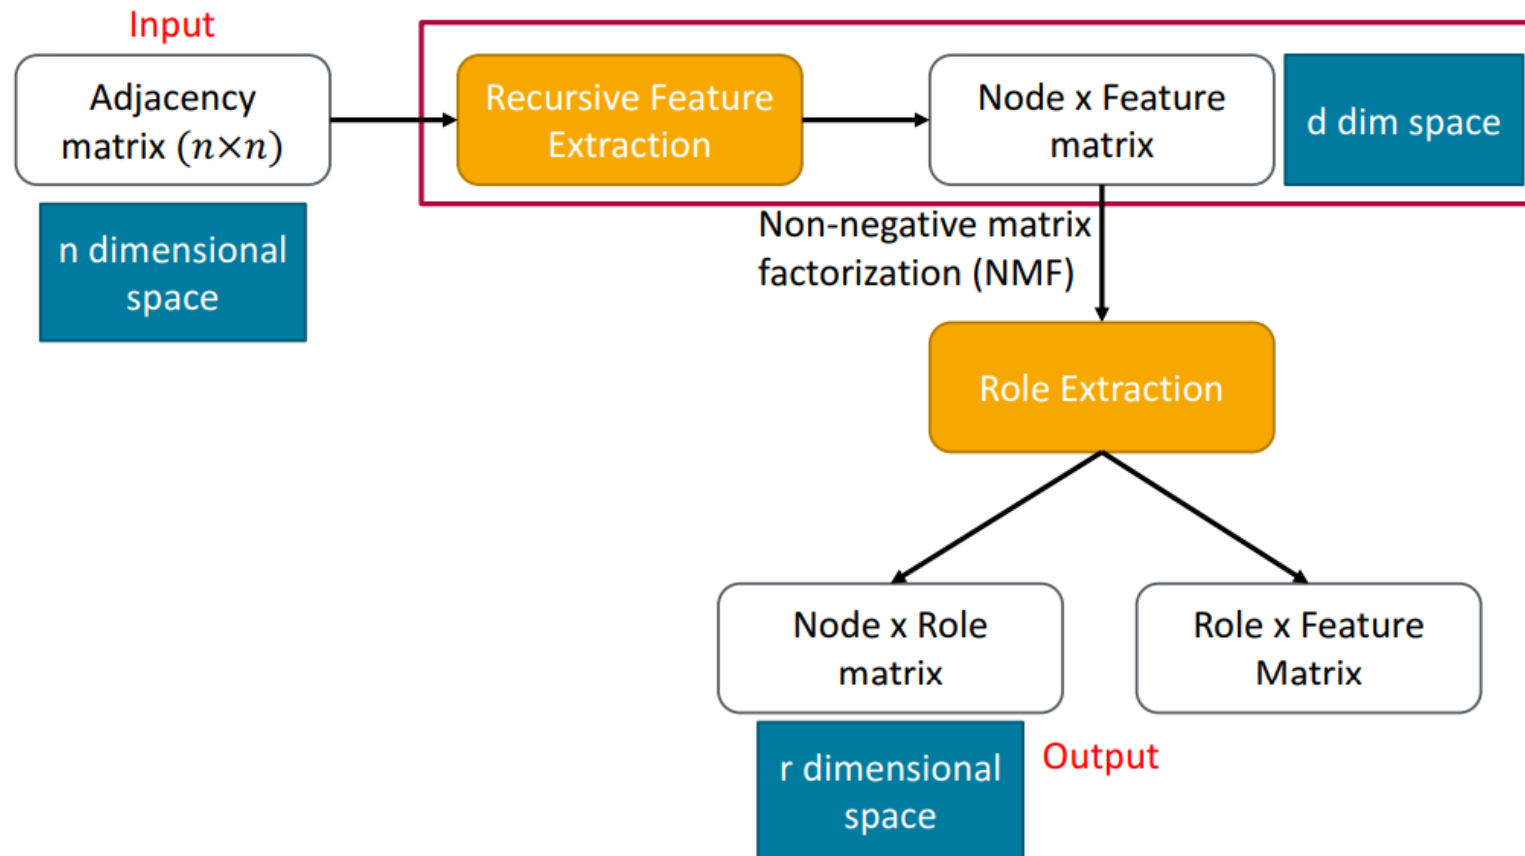
- - - - -> Collaborative-Based Filtering Recommendation



- What is the relation among the three equivalences?



- Takes features extracted with **ReFeX** and **factorizes the binary node-feature matrix** in order to create low dimensional structural node representations




```
# assign node roles
role_extractor = RoleExtractor(n_roles=None)
role_extractor.extract_role_factors(features)
node_roles = role_extractor.roles

print('\nNode role assignments:')
pprint(node_roles)


print('\nNode role membership by percentage:')
print(role_extractor.role_percentage.round(2))
```

Node role assignments:

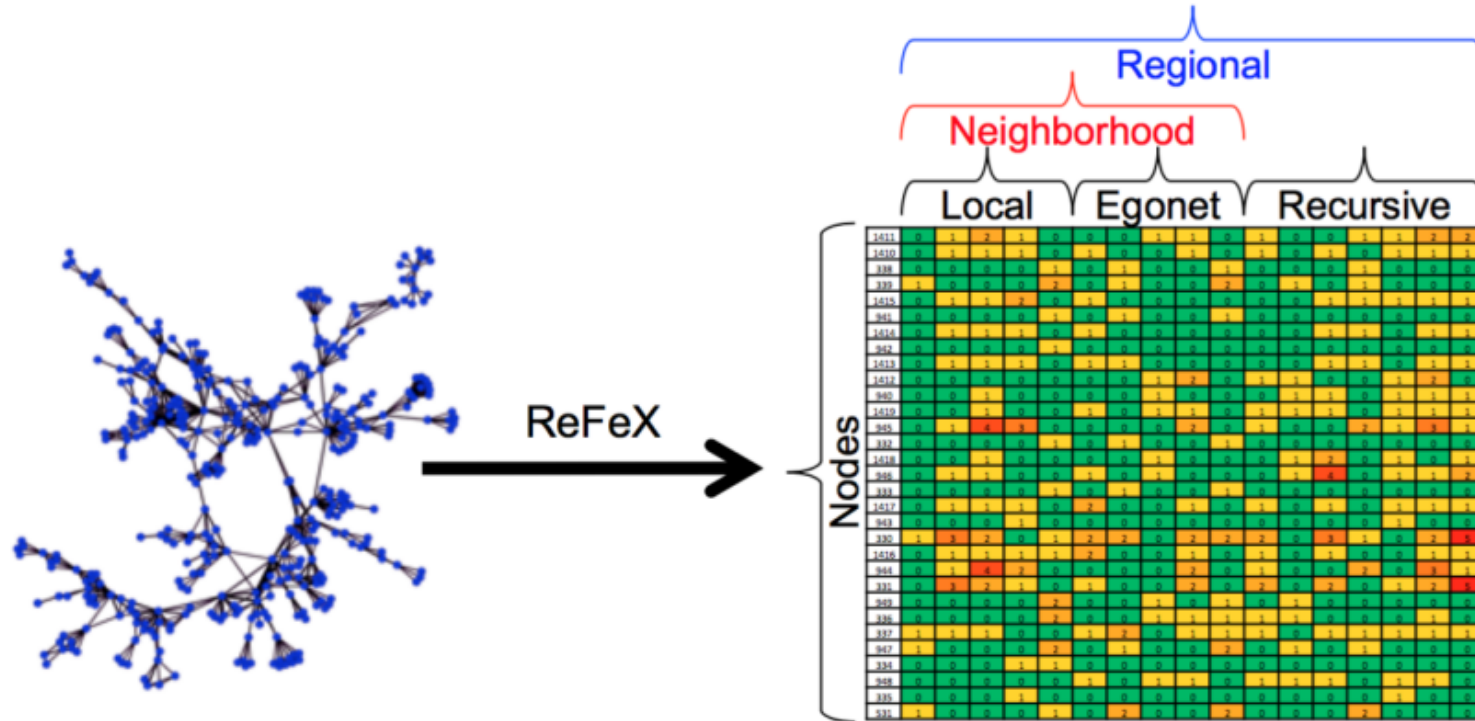
{0: 'role_1', 1: 'role_0', 2: 'role_1', 3: 'role_0', 4: 'role_1'}

Node role membership by percentage:

	role_0	role_1
0	0.03	0.97
1	0.97	0.03
2	0.25	0.75
3	0.69	0.31
4	0.25	0.75




- Transform the network **connectivity** into **recursive structural features**.
- Technically, embeds the graph into an $|\mathcal{F}|$ dimensional space, where \mathcal{F} is a set of **features** (degree, self-loops, avg edge weight, # of edges in egonet)



- Local:
 - Measures of the node **degree**
- Egonet:
 - The **egonet** (or ego-network) of a node is the node itself, the adjacent nodes, and the graph induced by those nodes
 - Computed based on each node's ego network: #of within-egonet edges, #of edges entering & leaving the egonet
- **Recursive**
 - Some **aggregate** (mean, sum, max, min, etc.) of another feature over a node's neighbours
 - The aggregation can be computed over any real-valued feature, including other recursive features.

- Find r overlapping clusters in the feature space
 - Each node can have multiple roles at the same time
- Generate a rank r approximation of the node \times feature matrix V
- Use **non-negative matrix factorization**: $V \approx GF$



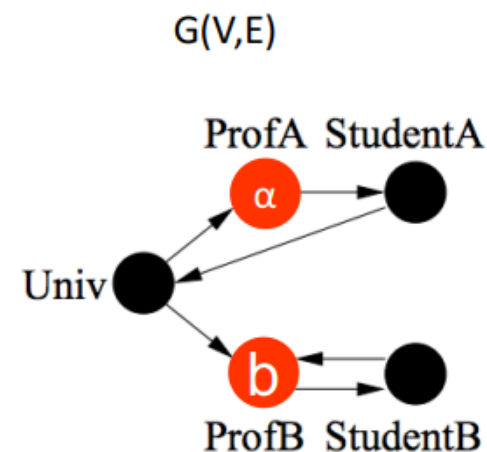
- The G matrix assigns nodes to roles
- The F matrix represents how the features explain the roles

- Idea:
 - Two objects are similar if they are **referenced by similar objects**

decay factor in $[0,1]$


$$s(a, b) = \frac{\text{Avg similarity between in-neighbors of } \alpha \text{ and in-neighbors of } b}{\text{total \# of in-neighbors pairs}}$$

similarity of in-neighbors

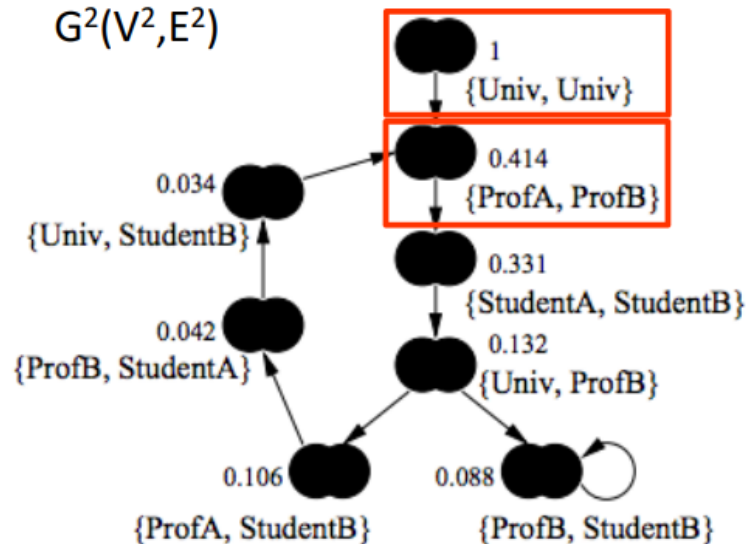


- Intuition: Computing SimRank is like **propagating on the G^2 graph of node-node pairs**
 - The **source** of similarity is **self-vertices**, like (Univ, Univ).
 - Similarity **propagates along pair-paths** in G^2 , away from the sources.

$G(V,E)$



$G^2(V^2,E^2)$




Structural context

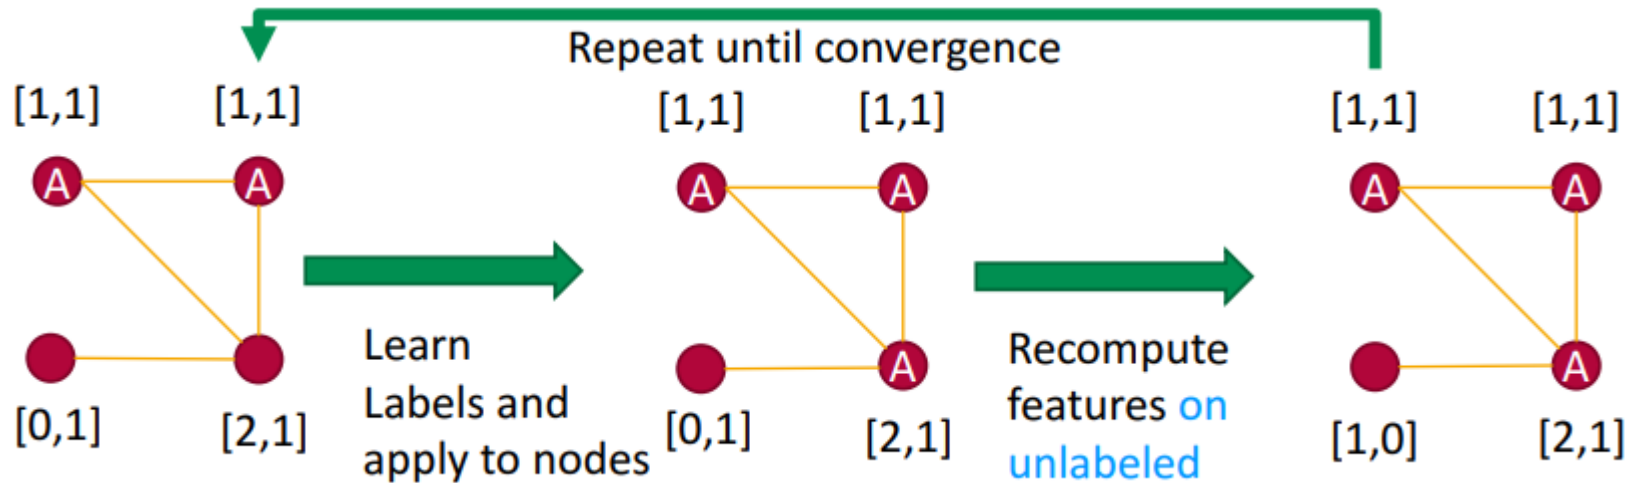
- Average similarity between A and B:

$$s(A, B) = \text{Avg similarity between out-neighbors of A and out-neighbors of B}$$

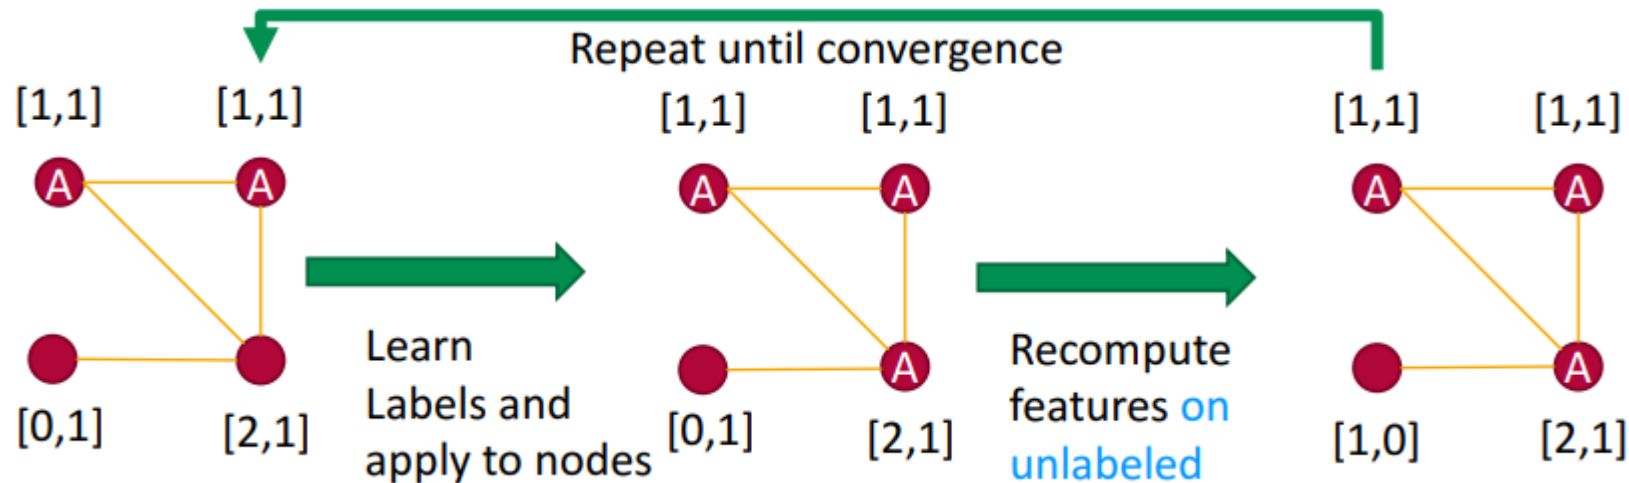
$$s(c, d) = \text{Avg similarity between in-neighbors of c and in-neighbors of d}$$



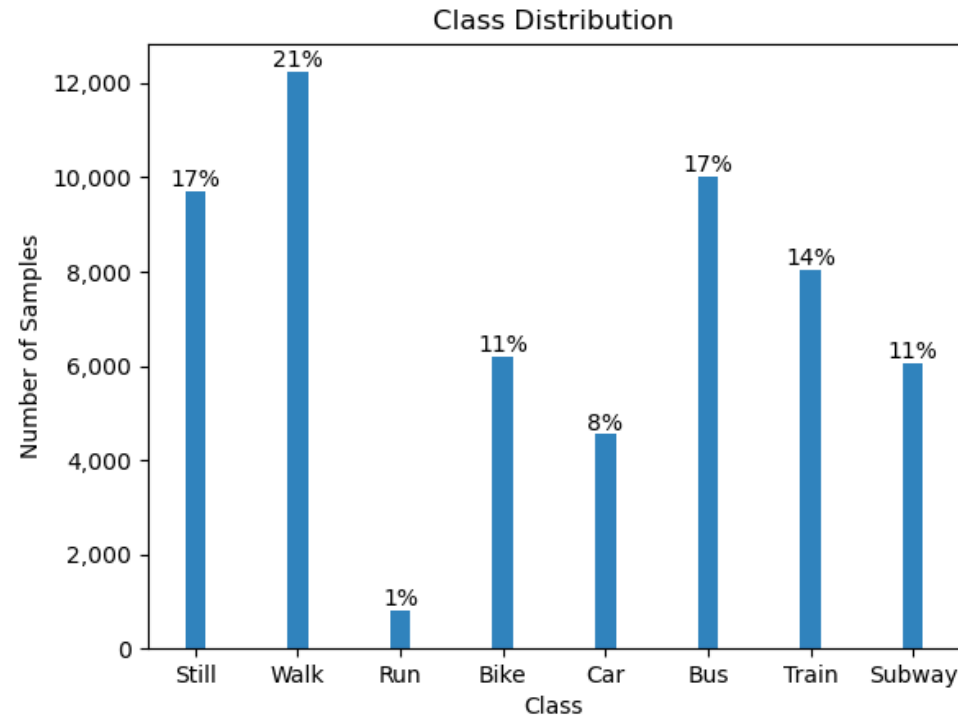
- Idea:
 - Use **features that consider the neighbor nodes**
 - and **repeat the classification** several time **until nothing changes**
- Suppose for each node we have two features:
 - Number of neighbors with class A
 - Number of neighbors without a class



- **Train classifier** using the labelled instances (SVM, Random Forest, etc.)
- Until convergence
 - **Apply classifier to the unlabelled nodes**
 - **Update the feature vectors** for unlabelled nodes
- Return the labels for the labelled nodes



- Each node has a **distribution over the labels**
- To avoid noise keep only the **top-k labels** for each unlabelled node sorted in descending order.
 - Intuition: **remove the less confident labels**



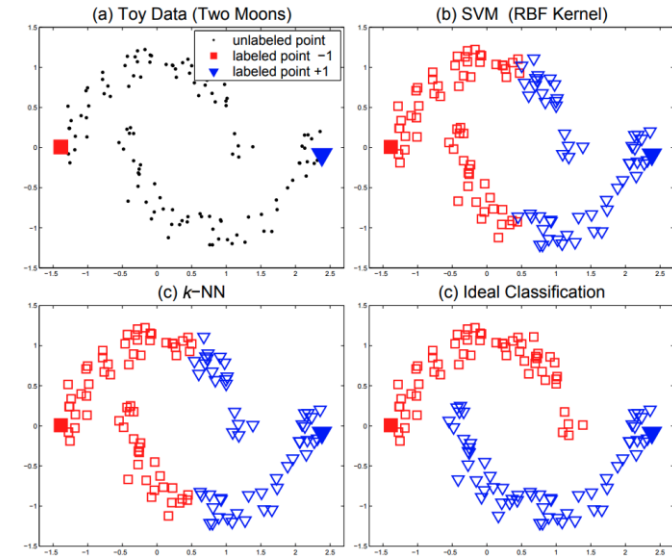
- The keynote of method is to let **every point iteratively spread its label information** to its neighbours **until a global stable state** is achieved.
 - Input:
 - Given a point set $X = \{x_1, \dots, x_l, x_{l+1}, \dots, x_n\} \in R^m$
 - A label set $L = \{1, \dots, c\}$, the first l points x_i are labelled as $y_i \in L$ and the remaining points x_u ($l + 1 \leq u \leq n$) are unlabelled.
 - Output:
 - Predict the label of the unlabelled points
-
- Let F denote the set of $n \times c$ matrices with nonnegative entries. A matrix $F = [F_1, F_2, \dots, F_n]$ corresponds to a classification on the set X by labelling each point x_i as a label $y_i = \operatorname{argmax}_{\{j \leq c\}} F_{ij}$.
 - We can understand F as a vectorial function $F: X \rightarrow R^c$ which assigns a vector F_i to each point x_i .


The algorithm is as follows:

1. Form the affinity matrix W defined by:

$$W_{ij} = \begin{cases} e^{-\frac{\|x_i - x_j\|^2}{2\sigma^2}}, & \text{if } i \neq j \\ 0 & , \text{ if } i = j \end{cases}$$



2. Construct the matrix $S = D^{-1/2}WD^{-1/2}$ where D is a diagonal matrix with its (i,i) - element equal to the **sum of the i -th row of W** .
3. Iterate $F(t + 1) = \alpha SF(t) + (1 - \alpha)Y$ until convergence, where $\alpha \in (0,1)$
3. Finally, the label of each unlabelled point is set to be the class of which it has received most information during the iteration process.





- denotes +1
- denotes -1

Support Vectors are those datapoints that the margin pushes up against



$$f(\mathbf{x}, \mathbf{w}, b) = \text{sign}(\mathbf{w} \cdot \mathbf{x} + b)$$

The **maximum margin linear classifier** is the linear classifier with the, um, maximum margin.

This is the simplest kind of SVM (Called an LSVM)

Linear SVM

```
from networkx.algorithms import node_classification
G = nx.path_graph(4)
G.nodes[0]['label'] = 'A'
G.nodes[3]['label'] = 'B'
G.nodes(data=True)


G.edges()

predicted = node_classification.local_and_global_consistency(G)
predicted

['A', 'A', 'B', 'B']
```

- **Random Walk**
- Personalized Random Walk with Restarts (RWR)
- PageRank: A kind of random walk
- Set of neighbours of nodes

- In a random walk, you assume that the **walker moves randomly** and **chooses one of the neighbours** to visit.
- In the figure:
 - A chooses B or C with probability $1/2$.
 - Once he chooses one it increases the number of times he visited that node
 - Continue the process until nothing changes anymore (at a probabilistic level)



D will receive many visits since many nodes are connected to him


- Now assume that with **probability c** you **perform another move** and with **probability $(1-c)$** you **jump back** to Tom.
- Therefore the probability for the walker of being in Tom place will be:

$$Tom(t) = c \frac{Prob(Andrew)(t-1) + Prob(Laura)(t-1)}{3} + (1-c)$$


Probability of visiting Andrew at time $(t-1)$

Number of Andrew and Laura's neighbors

Probability of jumping back to Tom



- Start **two random walks from the two nodes** you want to compare separately
- **Compare the final scores** you obtain for each node in the graph using some vector comparison (e.g., cosine similarity, KL-divergence)



[Tom, Andrew, Alice, John, ..., Paul]


Vector for Tom = [0.2, 0.2, 0.1, 0.3, ..., 0.01]

Vector for John = [0.05, 0.15, 0.2, 0.2, ..., 0.2]


Compare the two
vectors (e.g., subtract)

- At each step of the process, **each vertex updates its label to a new one** which corresponds to **the most frequent label among its neighbours**.
- For each vertex $v \in V$, v updates its label according to:


$$l_v = \arg \max_l \sum_{u \in N(v)} [l_u == l]$$




(a) Result after the first label propagation



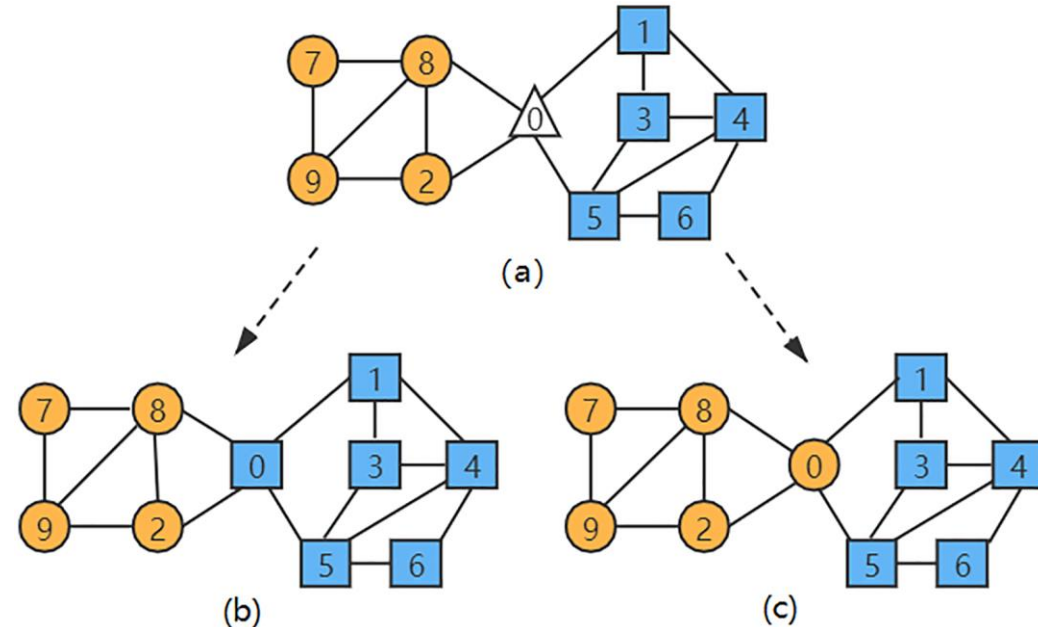
(b) Result after the second label propagation




(c) Result after the third label propagation



(d) Result after the fourth label propagation





```
colors = ["#00C98D", "#5030C0", "#50F0F0"]
pos = nx.spring_layout(G)
lst_m = community.label_propagation_communities(G)
color_map_b = {}
keys = G.nodes()
values = "black"
for i in keys:
    color_map_b[i] = values
counter = 0
for c in lst_m:
    for n in c:
        color_map_b[n] = colors[counter]
        counter = counter + 1
nx.draw_networkx_edges(G, pos)
nx.draw_networkx_nodes(G, pos, node_color=dict(color_map_b).values())
nx.draw_networkx_labels(G, pos)
plt.axis("off")
plt.show()
```



- Precision/Recall
- Accuracy + weighted loss
- ROC and AUC

- When evaluating a search tool or a classifier, we are interested in at least two performance measures:
- **Precision**: Within a given set of positively-labeled results, the fraction that were true positives = $tp/(tp + fp)$
- **Recall**: Given a set of positively-labeled results, the fraction of all positives that were retrieved = $tp/(tp + fn)$
- Positively-labeled means judged “relevant” by the search engine or labeled as in the class by a classifier. tp = true positive, fp = false positive etc.

- Search tools and classifiers normally assign **scores** to items. Sorting by score gives us a precision-recall plot which shows what performance would be for **different score thresholds**.



- The simplest measure of performance would be the fraction of items that are correctly classified, or the “accuracy” which is:

$$\frac{tp + tn}{tp + tn + fp + fn}$$

- But this measure is **dominated by the larger set (of positives or negatives)** and favours trivial classifiers.
- e.g. if 5% of items are truly positive, then a classifier that always says “negative” is 95% accurate.

We can instead try to minimize a **weight sum**:

$$w_1 \text{fn} + w_2 \text{fp}$$

And typically, $w_1 \gg w_2$, since positives are often much rarer (clicks or purchases or viewing a movie).

A measure that naturally combines precision and recall is the β -weighted F-measure:

$$F = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$


Which is the weighted **harmonic mean of precision and recall**. Setting $\beta = 1$ gives us the F_1 – measure. It can also be computed as:

$$F_{\beta=1} = \frac{2PR}{P + R}$$


ROC is **Receiver-Operating Characteristic**. ROC plots

Y-axis: true positive rate = $tp/(tp + fn)$, same as recall

X-axis: false positive rate = $fp/(fp + tn) = 1 - \text{specificity}$



- ROC AUC is the “**Area Under the Curve**” – a single number that captures the overall quality of the classifier. It should be between 0.5 (random classifier) and 1.0 (perfect).





네트워크 과학연구실
NETWORK SCIENCE LAB



가톨릭대학교
THE CATHOLIC UNIVERSITY OF KOREA

